

Bottleneck Attention: Hard-Wiring Low-Dimensional Routing in GPT-Style Transformers

Daniel Owen van Dommelen
theapemachine@gmail.com

December 2025

Abstract

Background: Adaptive Low-Rank Training (ALRT) indicates extreme redundancy in Transformer attention: query/key (Q/K) projections compress to effective ranks $\approx 9\text{--}14$ (mean ≈ 11.2) out of 512 in prior experiments, suggesting that attention pattern computation may live in a low-dimensional comparison subspace.

Methods: We test a minimal architectural hypothesis (*Bottleneck Attention*): decouple residual stream width d_{model} from the attention subspace dimension d_{attn} . We compute attention with $Q, K, V : \mathbb{R}^{d_{\text{model}}} \rightarrow \mathbb{R}^{d_{\text{attn}}}$ where $d_{\text{attn}} \ll d_{\text{model}}$, and project back with $\mathbb{R}^{d_{\text{attn}}} \rightarrow \mathbb{R}^{d_{\text{model}}}$. Optionally, we append a learnable per-head null key/value pair $(k_\emptyset, v_\emptyset)$ to the key/value sequence, allowing the attention distribution to assign explicit probability mass p_\emptyset to a no-op sink (*attend nowhere*), and we learn per-head temperature scaling.

Results: On a controlled WikiText-2 word-level setup (6-layer GPT-style LM; $d_{\text{model}} = 512$, 8 heads, context 256), reducing d_{attn} from 512 to 128 (4 \times) and to 32 (16 \times , with a null slot) yields only modest degradation in best validation loss ($5.3688 \rightarrow 5.4800 \rightarrow 5.6070$). Notably, during early training (< 500 steps), bottleneck variants achieve lower validation perplexity than the full-width baseline (e.g., step 200: 582.11 baseline vs 518.64 at $d_{\text{attn}} = 128$; 519.18 at $d_{\text{attn}} = 32+\text{null}$), suggesting accelerated initial learning on a constrained optimization manifold.

Conclusions: Attention routing in this regime is intrinsically low-dimensional. Because KV-cache memory scales linearly in d_{attn} , reducing d_{attn} from 512 to 128 reduces per-token inference state by 75% (and to 32 by 93.75%), directly improving long-context inference feasibility.

1 Introduction

Transformers remain architecturally conservative despite dramatic changes in scale, hardware, and use cases since their introduction [1]. Many design choices persist because they work, not because they are known to be minimal. ALRT was motivated by a simple suspicion: much of Transformer parameterization is redundant, and this redundancy is not uniform across components. Empirically, ALRT finds strong projection-type asymmetry: attention Q/K projections compress far more aggressively than attention output and MLP projections [9, 10].

This report tests the most literal architectural consequence: if

attention pattern computation truly lives in a small subspace, we can *hard-wire* that by decoupling attention’s internal dimension from the residual stream width. Unlike dynamic rank resizing (which complicates optimizer state validity), this is a static architectural change trained from scratch.

2 Background: ALRT and the low-rank attention signal

ALRT parameterizes linear maps as low-rank factors $W = UV$ and adapts rank during training using stable rank estimation:

$$r_{\text{stable}}(W) = \frac{\|W\|_F^2}{\sigma_{\max}(W)^2}.$$

Across GPT-style language modeling experiments, ALRT reports that Q/K projections are the most compressible, with mean ranks ≈ 11.2 out of 512, while attention output and MLP projections retain higher effective ranks [9]. A natural architectural interpretation is that attention’s *comparison space* (used to compute patterns) is low-dimensional, even if the residual stream is not.

3 Method: Bottleneck Attention

3.1 Decoupled attention subspace

We replace full-width Q/K/V projections with projections into a smaller attention subspace:

$$W_Q, W_K, W_V \in \mathbb{R}^{d_{\text{model}} \times d_{\text{attn}}}, \quad W_O \in \mathbb{R}^{d_{\text{attn}} \times d_{\text{model}}}.$$

Attention score and apply operations now scale in d_{attn} , while the residual stream stays at d_{model} . This directly targets attention FLOPs, attention parameters, and KV-cache memory.

3.2 “Attend nowhere” via a null key/value (optional)

Softmax attention forces each query to allocate probability mass across keys. We optionally append a learnable per-head null key/value $(k_\emptyset, v_\emptyset)$ to the Key/Value sequence. For a query position t , the attention distribution becomes $p(\cdot | t) = \text{softmax}(\dots, q_t^\top k_\emptyset)$, allowing explicit probability mass p_\emptyset on a no-op sink. We initialize $v_\emptyset = 0$ so “attend nowhere” starts as “contribute nothing”.

Table 1: WikiText-2 (word-level) results. Best validation metrics over training.

Model	d_{attn}	Params (M)	Best val loss	Best va
Baseline	512	36.06	5.3688	21
Bottleneck	128	31.34	5.4800	23
Extreme+Null	32	30.16	5.6070	27

3.3 Learned per-head temperature (optional)

When d_{attn} is small, head dimension $d_{\text{head}} = d_{\text{attn}}/n_{\text{head}}$ becomes tiny and dot products can become noisy. We include a learnable per-head logit scale multiplying the standard $1/\sqrt{d_{\text{head}}}$ factor.

4 Experimental Setup

Dataset. WikiText-2 word-level setup using a minimal whitespace tokenizer with per-line `<eos>` boundaries and `<unk>` for OOV. A single input file is tokenized and split 90/10 into train/validation within the script.

Model. Decoder-only GPT-style Transformer with pre-norm blocks; 6 layers; $d_{\text{model}} = 512$; 8 heads; $d_{\text{ff}} = 2048$; context length 256; dropout 0.1; weight tying between input embedding and LM head.

Optimization. AdamW with lr=3e-4, weight decay=0.01, betas=(0.9,0.95), batch size 32, 200 steps/epoch, 30 epochs (6000 steps), gradient clipping 1.0. Evaluation every 200 steps (50 eval iters). Hardware: Apple MPS backend.

5 Results

Table 1 summarizes best validation metrics. Reducing d_{attn} yields a smooth degradation curve rather than a catastrophic cliff.

5.1 Early convergence anomaly

Figure 1 shows validation perplexity vs training step. Notably, at the first evaluation point (step 200), the bottleneck variants outperform the baseline (e.g., baseline 582.11 vs 518.64 for $d_{\text{attn}} = 128$). This pattern persists through the first several hundred steps (inset), contradicting the intuition that wider attention must be retained to “find” a good solution.

5.2 Pareto curve

Figure 2 shows best validation loss vs d_{attn} (log scale). Figure 3 shows parameters vs best validation loss.

6 Discussion

Attention as low-dimensional routing. The viability of $d_{\text{attn}} = 32$ suggests that attention can route effectively using

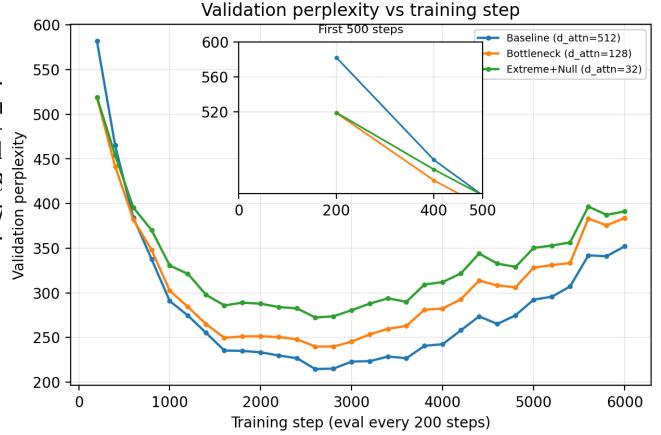


Figure 1: Validation perplexity vs training step (eval every 200 steps). Inset: first 500 steps. Bottleneck variants learn faster early despite lower attention dimensionality.

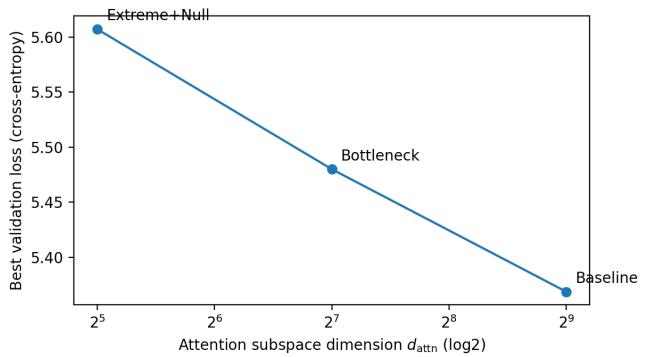


Figure 2: Best validation loss vs d_{attn} (log₂ x-axis).

a small comparison/mixing subspace, consistent with ALRT’s observation that Q/K matrices collapse to tiny effective rank. A useful mental model is “wide residual stream + narrow router.”

Why degradation is smooth. If high-dimensional attention geometry were essential for core competence, we would expect collapse once d_{attn} falls below a critical threshold. Instead, loss increases gradually, suggesting many attention patterns in this regime are representable with relatively few degrees of freedom.

Early learning signal. The early-phase perplexity advantage (Figure 1) suggests that full-width attention may be not only redundant but also *optimization-suboptimal* in this regime. One interpretation is that over-parameterized Q/K spaces introduce many weakly-aligned gradient directions, increasing noise and slowing early feature acquisition. Constraining the routing manifold may act as an implicit regularizer that accelerates early training.

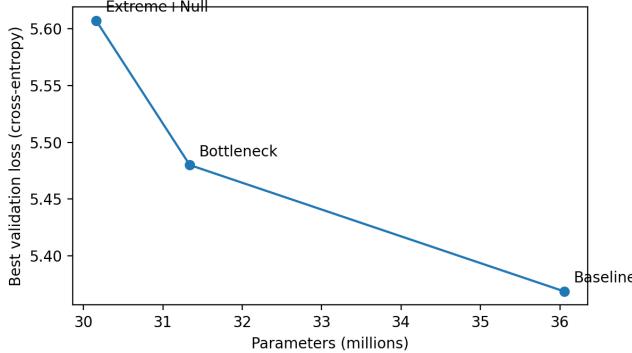


Figure 3: Parameters (M) vs best validation loss.

KV-cache memory: concrete scaling. For standard multi-head attention, per-token KV state per layer is

$$\text{KV per token per layer} = 2 \cdot n_{\text{head}} \cdot d_{\text{head}} = 2d_{\text{attn}}.$$

Therefore total KV-cache memory for batch size B , context length T , L layers, and element size b bytes is

$$M_{\text{KV}} = 2BTL d_{\text{attn}} b,$$

which scales *linearly* in d_{attn} . Reducing d_{attn} from 512 to 128 cuts KV-cache by 75% ($4\times$ smaller); reducing to 32 cuts KV-cache by 93.75% ($16\times$ smaller). In a hypothetical 128k-context regime, a $4\times$ reduction in per-token KV state enables $\approx 4\times$ larger batch size (or $\approx 4\times$ longer context) for the same memory budget, directly targeting a dominant modern inference bottleneck.

7 Related work

Low-rank adaptation and training. LoRA [2] demonstrates that many learned transformations can be modified via low-dimensional updates. ALRT extends this line of thought by measuring and adapting rank during training [9, 10]. AdaRankGrad [7] provides complementary theoretical and algorithmic evidence that learning signals themselves become low-rank: it proves that estimated layer gradients decrease in rank during training and asymptotically approach rank one, and leverages this structure to reduce optimizer memory. Unlike AdaRankGrad (which compresses *updates*), Bottleneck Attention applies low-dimensionality as a static architectural constraint, yielding inference-time KV-cache savings.

Efficient attention and KV-cache reduction. Efficient-attention methods address quadratic attention cost via sparsity or kernel/low-rank approximations (e.g., Reformer [3], Linformer [4], Performer [5]). Recent architectures such as DeepSeek-V2’s Multi-Head Latent Attention (MLA) compress the Key/Value representation into a low-rank latent to reduce inference memory [6]. Bottleneck Attention is complementary: it reduces the dimensionality of the *interaction space* (including the Query projection and the score/apply computation), not only the stored KV state.

8 Limitations and future work

Long-context scaling risk (Johnson–Lindenstrauss). These experiments use context length 256 and learned absolute positional embeddings. If attention must reliably discriminate among N candidate keys, random-projection theory suggests the dimensionality required to preserve pairwise distances scales as $O(\log N)$ (e.g., Johnson–Lindenstrauss-type bounds [8]). Thus $d_{\text{attn}} = 32$ may saturate at much longer sequence lengths, and the optimal d_{attn} may increase (slowly) with context length.

Position encoding. Interactions with RoPE/ALiBi and long-context regimes remain future work.

Variance and scale. We report single runs at modest scale; multi-seed sweeps and larger models are needed for confidence intervals and generalization.

9 Conclusion

ALRT indicates that attention pattern computation is intrinsically low-rank, with Q/K compressing to mean ranks $\approx 11.2/512$. Bottleneck Attention operationalizes this as a static architectural change: compute attention in a reduced subspace $d_{\text{attn}} \ll d_{\text{model}}$ while preserving a wide residual stream. On WikiText-2 (word-level), $d_{\text{attn}} = 128$ and even $d_{\text{attn}} = 32$ (with a null slot) remain viable with modest loss penalties, and show an early-phase perplexity advantage. Together, these results suggest that the standard convention $d_{\text{attn}} = d_{\text{model}}$ is substantial structural bloat in this regime.

Reproducibility

Code, scripts, and exact reproduction instructions (including pinned dependencies and commands) are provided at: <https://github.com/TheApeMachine/bottleneck-attention>

Conflict of interest statement

The author declares no competing interests.

Funding statement

This research received no external funding.

Ethics statement

This work does not involve studies with human participants, human data, human tissue, or animals; therefore, no ethical approval was required.

Data availability statement

The data supporting this study are derived from the publicly available WikiText-2 corpus. Training scripts, configuration, and logs sufficient to reproduce the reported results are available in the repository listed above.

References

- [1] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [2] Edward J Hu, Yelong Shen, Phil Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. *arXiv preprint arXiv:2106.09685*, 2021.
- [3] Nikita Kitaev, Łukasz Kaiser, and Anselm Levskaya. Reformer: The efficient transformer. In *International Conference on Learning Representations (ICLR)*, 2020.
- [4] Sinong Wang, Belinda Z Li, Madian Khabsa, Han Fang, and Hao Ma. Linformer: Self-attention with linear complexity. *arXiv preprint arXiv:2006.04768*, 2020.
- [5] Krzysztof Choromanski, Valerii Likhoshesterov, David Dohan, Xingyou Song, Andreea Gane, Tamas Sarlos, Peter Hawkins, Jared Davis, Afroz Mohiuddin, Łukasz Kaiser, and others. Rethinking attention with performer’s random features. In *International Conference on Learning Representations (ICLR)*, 2021.
- [6] DeepSeek-AI. DeepSeek-V2: A strong, economical, and efficient mixture-of-experts language model. *arXiv preprint arXiv:2405.04434*, 2024.
- [7] Yehonathan Refael, Jonathan Svirsky, Boris Shustein, Wasim Huleihel, and Ofir Lindenbaum. AdaRankGrad: Adaptive gradient-rank and moments for memory-efficient LLMs training and fine-tuning. *arXiv preprint arXiv:2410.17881*, 2024.
- [8] William B Johnson and Joram Lindenstrauss. Extensions of Lipschitz mappings into a Hilbert space. In *Conference in Modern Analysis and Probability*, pages 189–206. American Mathematical Society, 1984.
- [9] Daniel Owen van Dommelen. Adaptive low-rank training via spectral rank estimation: Dynamic parameter efficiency in transformer language models (technical report). Internal technical report, December 2025.
- [10] Daniel Owen van Dommelen. Adaptive low-rank training — technical report v2. Internal technical report, December 2025.